



L26- VIRTUAL MEMORY

Overview of Memory Management

- ❖ Demand Paging
- ❖ Copy-on-Write
- ❖ Page Replacement
- ❖ Allocation of Frames
- ❖ Thrashing
- ❖ Memory-Mapped Files

Objectives

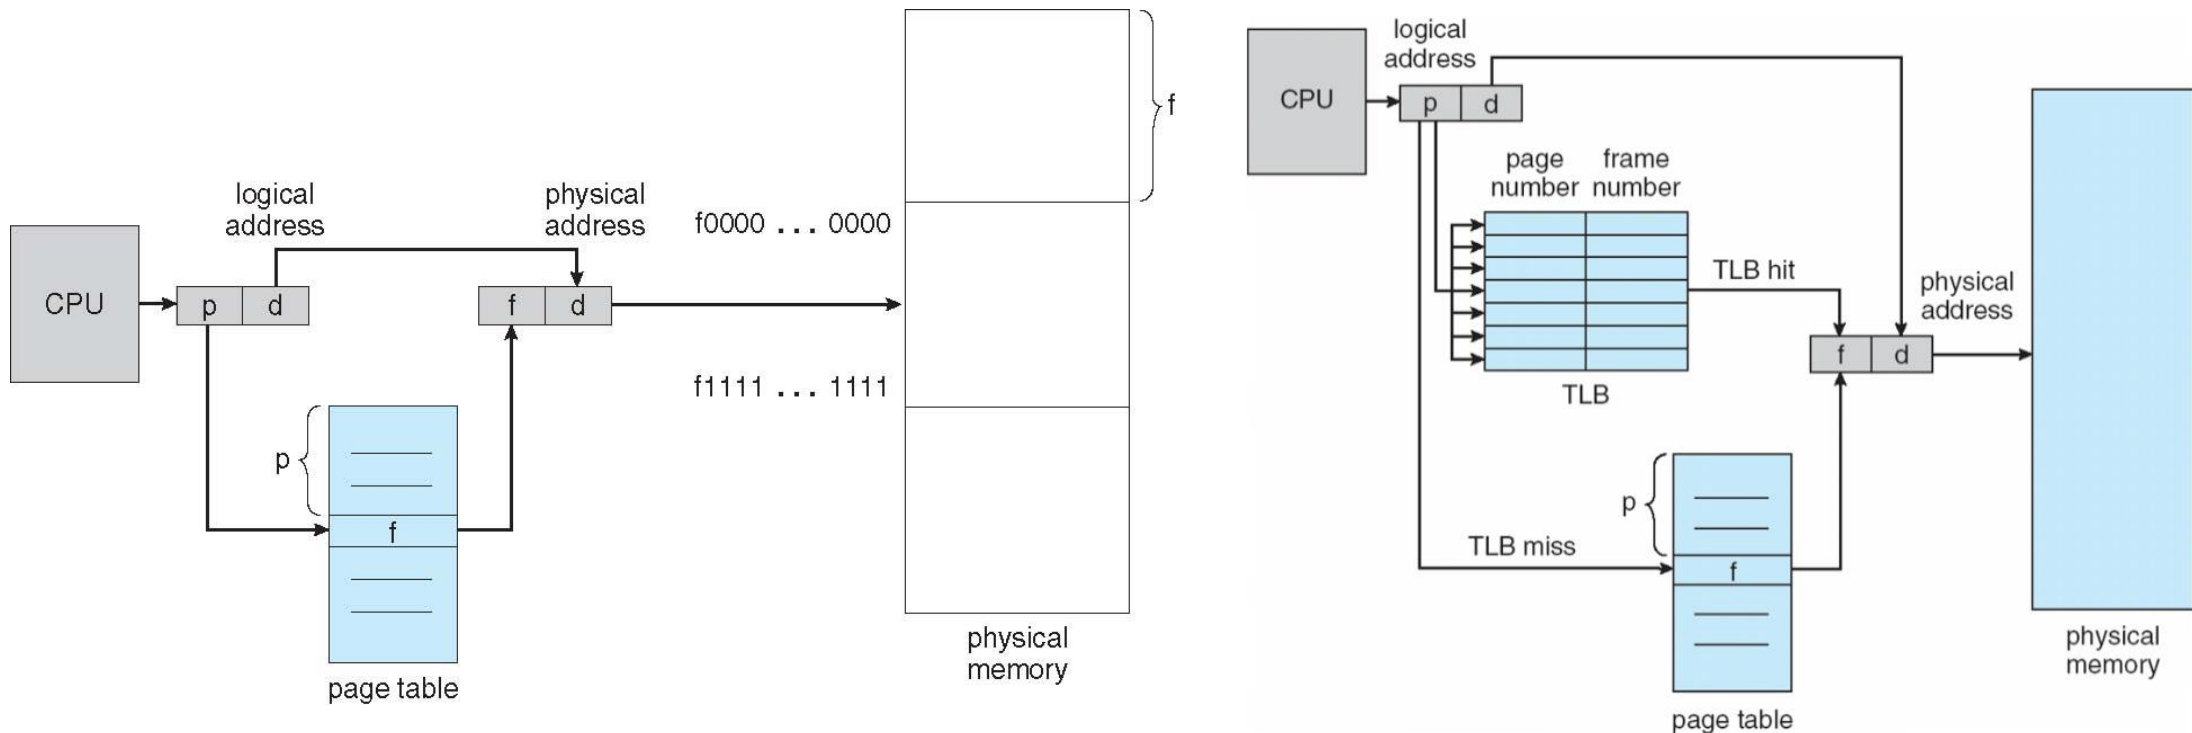
- ❖ To describe the benefits of a virtual memory system
- ❖ To explain the concepts of demand paging, page-replacement algorithms, and allocation of page frames
- ❖ To discuss the principle of the working-set model
- ❖ To examine the relationship between shared memory and memory-mapped files

Paging & Address Translation Scheme

❖ Address generated by CPU is divided into:

- ❖ **Page number** (*p*) – used as an index into a page table
- ❖ **Page offset** (*d*) – displacement within a page

page number	page offset
p	d
m - n	n



Structure of the Page Table

- ❖ Memory structures for paging can get huge using straight-forward methods
 - ❖ Consider a 32-bit logical address space as on modern computers
 - ❖ Page size of 4 KB (2^{12})
 - ❖ Page table would have 1 million entries ($2^{32} / 2^{12}$)
 - ❖ If each entry is 4 bytes -> 4 MB of physical address space / memory for page table alone
- ❖ Hierarchical Paging
- ❖ Hashed Page Tables
- ❖ Inverted Page Tables

Background

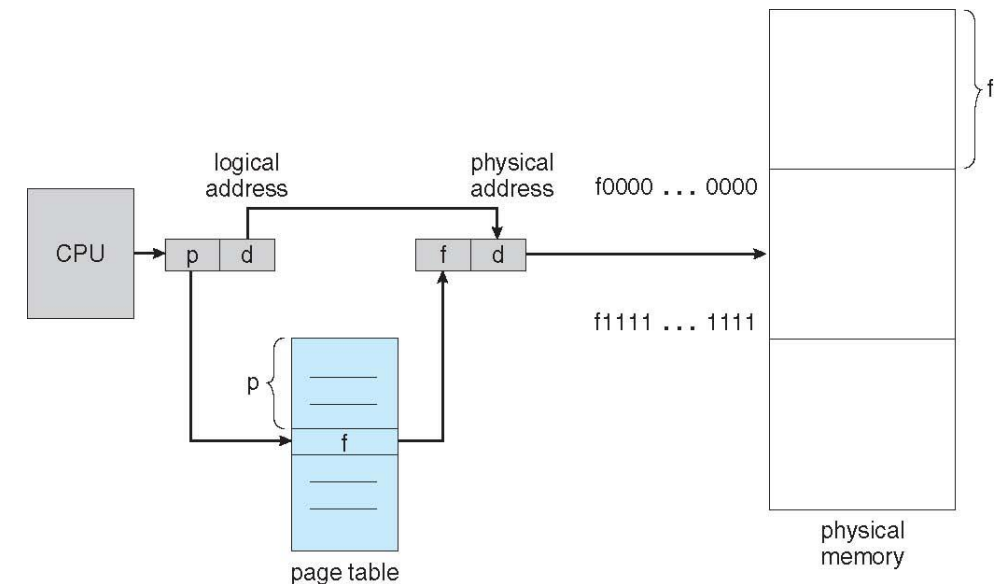
- ❖ Code needs to be in memory to execute, but entire program rarely used
 - ❖ Error code, unusual routines, large data structures
- ❖ Entire program code not needed at same time
- ❖ Consider ability to execute partially-loaded program
 - ❖ Program no longer constrained by limits of physical memory
 - ❖ Each program takes less memory while running → more programs run at the same time
 - ❖ Increased CPU utilization and throughput with no increase in response time or turnaround time

Background

- ❖ **Virtual memory** – separation of user logical memory from physical memory
- ❖ Only part of the program needs to be in memory for execution
- ❖ Logical address space can therefore be much larger than physical address space
- ❖ Allows address spaces to be shared by several processes
- ❖ More programs running concurrently
- ❖ Less I/O needed to load or swap processes


Background

- ❖ **Virtual address space** – logical view of how process is stored in memory
 - ❖ Usually start at address 0, contiguous addresses until end of space
 - ❖ Meanwhile, physical memory organized in page frames
 - ❖ **MMU** must map logical to physical
- ❖ Virtual memory can be implemented via:
 - ❖ Demand paging
 - ❖ Demand segmentation




Virtual-address Space

- ❖ Logical address space for stack to start at Max logical address and grow “down” while heap grows “up”
- ❖ Unused address space between stack and heap is hole
- ❖ Enables **sparse** address spaces with holes left for growth, dynamically linked libraries, etc




Virtual Memory To Physical Memory Mapping



Demand Paging

- ❖ Bring a page into memory only when it is needed
 - ❖ Less I/O needed, no unnecessary I/O
 - ❖ Less memory needed
 - ❖ Faster response
 - ❖ More users
- ❖ **Lazy swapper** – never swaps a page into memory unless page will be needed
 - ❖ Swapper that deals with pages is a **pager**



Valid-Invalid Bit

- ❖ With each page table entry a valid–invalid bit is associated (**v** \Rightarrow in-memory – **memory resident**, **i** \Rightarrow not-in-memory)
- ❖ Initially valid–invalid bit is set to **i** on all entries

Frame #	valid-invalid bit
	v
	v
	v
	i
...	
	i
	i

page table

0	A
1	B
2	C
3	D
4	E
5	F
6	G
7	H


logical memory

frame	valid-invalid bit
0	4 v
1	i
2	6 v
3	i
4	i
5	9 v
6	i
7	i

page table

0
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15

physical memory




Page Table When Some Pages Are Not in Main Memory

Page Fault

- ❖ If there is a reference to a page, first reference to that page will trap to operating system: **page fault** (page not found in main memory)
- ❖ Find free frame
- ❖ Swap page into frame via scheduled disk operation
- ❖ Reset tables to indicate page now in memory: Set validation bit = **1**
- ❖ Restart the instruction that caused the page fault

Steps in Handling a Page Fault



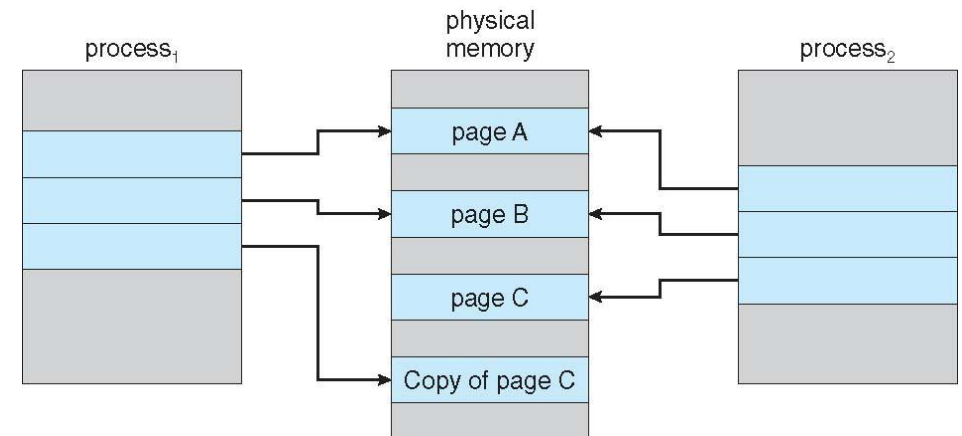

Demand Paging Overhead

- ❖ Page Fault Rate $0 \leq p \leq 1$
 - if $p = 0$ no page faults
 - if $p = 1$, every reference is a fault
- ❖ Effective Access Time (EAT)

$$\text{EAT} = (1 - p) \times \text{memory access} + p (\text{page fault overhead} \\ + \text{swap page out} + \text{swap page in})$$

Copy-on-Write


- ❖ **Copy-on-Write** (COW) allows both parent and child processes to initially **share** the same pages in memory
- ❖ If either process modifies a shared page, only then is the page copied
- ❖ COW allows more efficient process creation as only modified pages are copied
- ❖ In general, free pages are allocated from a **pool** of **zero-fill-on-demand** pages



Basic Page Replacement

- ❖ Find the location of the desired page on disk
- ❖ Find a free frame:
 - If there is a free frame, use it
 - If there is no free frame, use a page replacement algorithm to select a **victim frame**
 - Write victim frame to disk if **dirty**
- ❖ Bring the desired page into the (newly) free frame; update the page and frame tables
- ❖ Continue the process by restarting the instruction that caused the trap

Page Replacement




Page Replacement Algorithm

❖ Page-replacement algorithm

- ❖ Want lowest page-fault rate on both first access and re-access
- ❖ Evaluate algorithm by running it on a particular string of memory references (reference string) and computing the number of page faults on that string
 - ❖ String is just page numbers, not full addresses
 - ❖ Repeated access to the same page does not cause a page fault
 - ❖ **FIFO, LIFO,**
 - ❖ **Optimal,**
 - ❖ **LRU, LRU approximations,**
 - ❖ **LFU, MFU**

Page Faults Vs Number of Frames



First-In-First-Out (FIFO) Algorithm

- ❖ Ref. string: **7,0,1,2,0,3,0,4,2,3,0,3,0,3,2,1,2,0,1,7,0,1**
- ❖ 3 frames (3 pages can be in memory at a time per process)

reference string

7 0 1 2 0 3 0 4 2 3 0 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2	2	2	4	4	4	0	0	0	0	0	0	0	0	7	7	7
	0	0	0	3	3	3	2	2	2	1	1	1	1	1	1	1	1	0	0
		1	1	1	0	0	0	3	3	3	2	2	2	2	2	2	2	2	1


page frames

15 page faults

- ❖ How to track ages of pages? - Use a FIFO queue

Belady's Anomaly

- ❖ Consider 1,2,3,4,1,2,5,1,2,3,4,5
- ❖ Adding more frames can cause more page faults! - Belady's Anomaly



Optimal Algorithm

- ❖ Replace page that will not be used for longest period of time
 - ❖ 9 is optimal for the example
- ❖ Practical difficulty- Can't read the future
- ❖ Used for measuring how well your algorithm performs

reference string


7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		2		2		2						7		
	0	0	0		0		0		0		0						0		
		1	1		3		3		3		1						1		

page frames

Least Recently Used (LRU) Algorithm

- ❖ Use past knowledge rather than future
- ❖ Replace page that has not been used in the most amount of time
- ❖ Associate time of last use with each page



- ❖ 12 faults – better than FIFO but worse than OPT
- ❖ Generally good algorithm and frequently used

LRU Algorithm Implementation

❖ Counter implementation

- ❖ Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter
- ❖ When a page needs to be changed, look at the counters to find smallest value.

- ❖ To implement the LRU algorithm, the memory controller must track the LRU block as the computation proceeds.
- ❖ Example: Consider a Page Table with 4 pages.
 - ❖ For tracking the LRU block within a Page table, we use a 2-bit counter with every block.
 - ❖ When hit occurs:
 - ❖ Counter of the referenced block is reset to 0.
 - ❖ Counters with values originally lower than the referenced one are incremented by 1, and all others remain unchanged.

- When miss occurs:
 - If the set is not full, the counter associated with the new block loaded is set to 0, and all other counters are incremented by 1.
 - If the set is full, the block with counter value 3 is removed, the new block put in its place, and the counter set to 0. The other three counters are incremented by 1.

<div> <div>x</div> <div>Block 0</div> </div> <div> <div>x</div> <div>Block 1</div> </div> <div> <div>x</div> <div>Block 2</div> </div> <div> <div>x</div> <div>Block 3</div> </div> <div>Initial</div>	<div> <div>x</div> <div>Block 0</div> </div> <div> <div>x</div> <div>Block 1</div> </div> <div> <div>0</div> <div>Block 2</div> </div> <div> <div>x</div> <div>Block 3</div> </div> <div>Miss: Block 2</div>	<div> <div>0</div> <div>Block 0</div> </div> <div> <div>x</div> <div>Block 1</div> </div> <div> <div>1</div> <div>Block 2</div> </div> <div> <div>x</div> <div>Block 3</div> </div> <div>Miss: Block 0</div>	<div> <div>1</div> <div>Block 0</div> </div> <div> <div>x</div> <div>Block 1</div> </div> <div> <div>2</div> <div>Block 2</div> </div> <div> <div>0</div> <div>Block 3</div> </div> <div>Miss: Block 3</div>	<div> <div>2</div> <div>Block 0</div> </div> <div> <div>0</div> <div>Block 1</div> </div> <div> <div>3</div> <div>Block 2</div> </div> <div> <div>1</div> <div>Block 3</div> </div> <div>Miss: Block 1</div>	<div> <div>0</div> <div>Block 0</div> </div> <div> <div>1</div> <div>Block 1</div> </div> <div> <div>3</div> <div>Block 2</div> </div> <div> <div>2</div> <div>Block 3</div> </div> <div>Hit: Block 0</div>
<div> <div>1</div> <div>Block 0</div> </div> <div> <div>2</div> <div>Block 1</div> </div> <div> <div>0</div> <div>Block 2</div> </div> <div> <div>3</div> <div>Block 3</div> </div> <div>Miss: Block 2</div>	<div> <div>2</div> <div>Block 0</div> </div> <div> <div>3</div> <div>Block 1</div> </div> <div> <div>1</div> <div>Block 2</div> </div> <div> <div>0</div> <div>Block 3</div> </div> <div>Hit: Block 3</div>	<div> <div>2</div> <div>Block 0</div> </div> <div> <div>3</div> <div>Block 1</div> </div> <div> <div>0</div> <div>Block 2</div> </div> <div> <div>1</div> <div>Block 3</div> </div> <div>Hit: Block 2</div>	<div> <div>0</div> <div>Block 0</div> </div> <div> <div>3</div> <div>Block 1</div> </div> <div> <div>1</div> <div>Block 2</div> </div> <div> <div>2</div> <div>Block 3</div> </div> <div>Hit: Block 0</div>	<div> <div>1</div> <div>Block 0</div> </div> <div> <div>0</div> <div>Block 1</div> </div> <div> <div>2</div> <div>Block 2</div> </div> <div> <div>3</div> <div>Block 3</div> </div> <div>Miss: Block 1</div>	<div> <div>1</div> <div>Block 0</div> </div> <div> <div>0</div> <div>Block 1</div> </div> <div> <div>2</div> <div>Block 2</div> </div> <div> <div>3</div> <div>Block 3</div> </div> <div>Hit: Block 1</div>

LRU Approximation Algorithms


❖ Reference bit

- ❖ With each page associate a bit, initially = 0
- ❖ When page is referenced, bit set to 1
- ❖ Replace any with reference bit = 0 (if one exists)


❖ Second-chance algorithm

- ❖ If page to be replaced has
 - ❖ Reference bit = 0 → replace it
 - ❖ reference bit = 1 then:
 - ❖ set reference bit 0, leave page in memory
 - ❖ replace next page, subject to same rules

Second-Chance Page-Replacement



(a)



(b)

Enhanced Second-Chance Algorithm

- ❖ Improve algorithm by using reference bit and modify bit
- ❖ Take ordered pair (reference, modify)
- ❖ (0, 0) neither recently used nor modified – best page to replace
- ❖ (0, 1) not recently used but modified – not quite as good, must write out before replacement
- ❖ (1, 0) recently used but clean – probably will be used again soon
- ❖ (1, 1) recently used and modified – probably will be used again soon and need to write out before replacement

Counting Algorithms

- ❖ Keep a counter of the number of references that have been made to each
- ❖ **Least Frequently Used (LFU) Algorithm**: replaces page with smallest count
- ❖ **Most Frequently Used (MFU) Algorithm**: based on the argument that the page with the smallest count was probably just brought in and has yet to be used

Page-Buffering Algorithms

- ❖ Always keep a pool of free frames
 - ❖ When needed, frame is always available, not found at fault time
 - ❖ Read page into free frame and select victim to evict and add to free pool
 - ❖ When convenient, evict victim
- ❖ Keep list of modified pages
 - ❖ When free, write to backing store and set to non-dirty
- ❖ Possibly, keep free frame contents intact and note what is in them
 - ❖ If referenced again before reused, no need to load contents again from disk
 - ❖ Generally useful to reduce penalty if wrong victim frame selected



Thank You