



L21-INTRODUCTION TO DEADLOCKS

Session Outline

- ❖ **System Model**
- ❖ **Deadlock Characterization**
- ❖ **Resource Allocation Graph**
- ❖ **Methods for Handling Deadlocks**
- ❖ **Deadlock Prevention**

Objectives of Deadlock Management Unit

- ❖ To develop a description of deadlocks, which prevent sets of concurrent processes from completing their tasks
- ❖ To present a number of different methods for preventing or avoiding deadlocks in a computer system

System Model

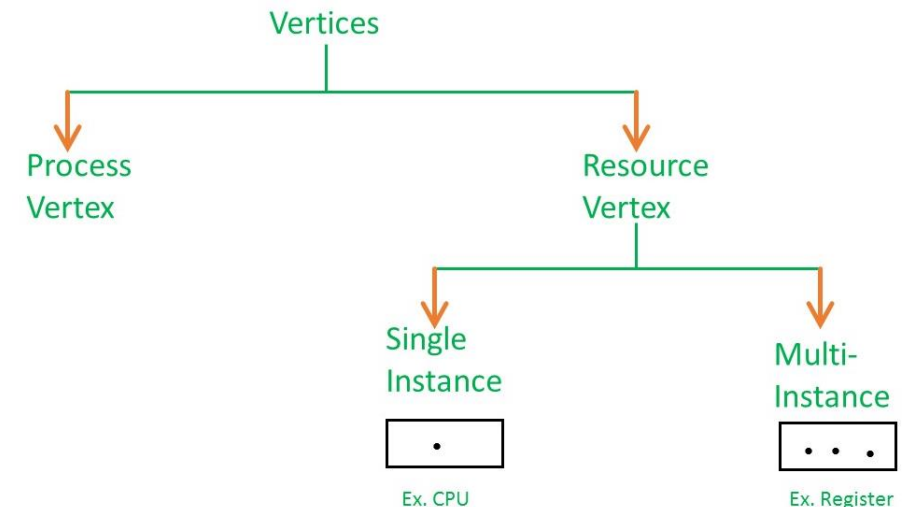
- ❖ System consists of resources
- ❖ Resource types R_1, R_2, \dots, R_m
 - ❖ *CPU cycles, memory space, I/O devices*
- ❖ Each resource type R_i has W_i instances.
- ❖ Each process utilizes a resource as follows:
 - ❖ **request**
 - ❖ **use**
 - ❖ **release**

Deadlock Characterization

- ❖ Deadlock can arise if the following four conditions hold simultaneously.
- ❖ **Mutual exclusion:** Only one process at a time can use a resource
- ❖ **Hold and wait:** A process holding at least one resource is waiting to acquire additional resources held by other processes
- ❖ **No preemption:** A resource can be released only voluntarily by the process holding it, after that process has completed its task
- ❖ **Circular wait:** There exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that P_0 is waiting for a resource that is held by P_1 , P_1 is waiting for a resource that is held by P_2 , ..., P_{n-1} is waiting for a resource that is held by P_n , and P_n is waiting for a resource that is held by P_0 .


Resource-Allocation Graph

- ❖ A set of vertices V and a set of edges E .
- ❖ V is partitioned into two types:
 - ❖ $P = \{P_1, P_2, \dots, P_n\}$, the set consisting of all the active processes in the system
 - ❖ $R = \{R_1, R_2, \dots, R_m\}$, the set consisting of all resource types in the system
- ❖ **request edge** – directed edge $P_i \rightarrow R_j$
- ❖ **assignment edge** – directed edge $R_j \rightarrow P_i$




Resource-Allocation Graph


❖ Process





❖ Resource Type with 4 instances





❖ P_i requests an instance of R_j




❖ P_i is holding an instance of R_j



Resource-Allocation Graph



RAG with a deadlock




RAG without a deadlock


Deadlock detection in RAG

- ❖ If graph contains no cycles \Rightarrow no deadlock
- ❖ If graph contains a cycle \Rightarrow
 - ❖ if only one instance per resource type, then deadlock
 - ❖ if several instances per resource type, possibility of deadlock

Deadlock detection in RAG



**A cycle...and
deadlock!**



**Same cycle...but no
deadlock. Why?**

Methods for Handling Deadlocks

- ❖ Ignore the problem and pretend that deadlocks never occur in the system; used by most operating systems, including UNIX
- ❖ Ensure that the system will **never** enter a deadlock state:
 - ❖ Deadlock prevention
 - ❖ Deadlock avoidance
- ❖ Allow the system to enter a deadlock state and then recover

Deadlock Prevention

- ❖ **Deadlock prevention is done by ensuring that at least one of the necessary 4 conditions for deadlock is not met.**
- ❖ **Mutual Exclusion** – not required for sharable resources (e.g., read-only files); must hold for non-sharable resources
- ❖ **Hold and Wait** – must guarantee that whenever a process requests a resource, it does not hold any other resources
 - ❖ Require process to request and be allocated all its resources before it begins execution, or allow process to request resources only when the process has none allocated to it.
 - ❖ Low resource utilization; starvation possible

Deadlock Prevention

❖ No Preemption –

- ❖ If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then all resources currently being held are released
 - ❖ Preempted resources are added to the list of resources for which the process is waiting
 - ❖ Process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting
- ## ❖ Circular Wait –
- impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration

Deadlock Example

```
/* thread one runs in this function */  
  
void *do_work_one(void *param)  
{  
  
    pthread_mutex_lock(&first_mutex);  
    pthread_mutex_lock(&second_mutex);  
  
    /** * Do some work */  
  
    pthread_mutex_unlock(&second_mutex);  
    pthread_mutex_unlock(&first_mutex);  
    pthread_exit(0);  
}
```

```
/* thread two runs in this function */  
  
void *do_work_two(void *param)  
{  
  
    pthread_mutex_lock(&second_mutex);  
    pthread_mutex_lock(&first_mutex);  
  
    /** * Do some work */  
  
    pthread_mutex_unlock(&first_mutex);  
    pthread_mutex_unlock(&second_mutex);  
    pthread_exit(0);  
}
```



Thank You